

Glow: Graph Lowering Compiler Techniques for Neural Networks

Nadav Rotem, Jordan Fix, Saleem Abdulrasool, Garret Catron, Summer Deng,
Roman Dzhabarov, Nick Gibson, James Hegeman, Meghan Lele, Roman Levenstein,
Jack Montgomery, Bert Maher, Satish Nadathur, Jakob Olesen, Jongsoo Park,
Artem Rakhov, Misha Smelyanskiy, Man Wang
Facebook

Abstract

This paper presents the design of Glow, a machine learning compiler for heterogeneous hardware. It is a pragmatic approach to compilation that enables the generation of highly optimized code for multiple targets. Glow lowers the traditional neural network dataflow graph into a two-phase strongly-typed intermediate representation. The high-level intermediate representation allows the optimizer to perform domain-specific optimizations. The lower-level instruction-based address-only intermediate representation allows the compiler to perform memory-related optimizations, such as instruction scheduling, static memory allocation and copy elimination. At the lowest level, the optimizer performs machine-specific code generation to take advantage of specialized hardware features. Glow features a lowering phase which enables the compiler to support a high number of input operators as well as a large number of hardware targets by eliminating the need to implement all operators on all targets. The lowering phase is designed to reduce the input space and allow new hardware backends to focus on a small number of linear algebra primitives.

1 Introduction

The end of power saving due to Moore’s Law, combined with the increased demand for compute power driven by machine learning, has led to a wave of innovation in computer architecture. Hennessy and Patterson [1] present five principles that guide the design of machine-learning domain specific architectures (DSA): dedicated local memories, large numbers of arithmetic units, simple forms of parallelism, reduced bit-widths, and domain-specific programming models. Compilers need to perform advance whole-graph

optimizations in order to execute neural networks efficiently on DSAs. This paper describes some of these techniques as implemented in Glow, an open-source machine learning compiler framework for heterogeneous hardware.

Traditional machine learning frameworks iterate over the nodes in the graph and execute them one by one. Unfortunately this node-visitor method of execution is inefficient, even on traditional processors. As a result, machine learning frameworks have started to hand over the graph to compilers [2, 3] that execute code more efficiently. Based on the increasing importance of neural networks, the need for energy efficiency in data centers and mobile devices, and the design principles of domain-specific architectures, we believe that the machine learning frameworks of the future will focus on providing attractive programming models on top of a layer that integrates compilers for many different targets.

In Glow, we focus on the lower parts of the software stack. We work to provide PyTorch [4] and other frameworks with a low-level graph and a code generator for neural networks. The name Glow is an abbreviation for Graph-Lowering, which is the main technique that the compiler uses for generating efficient code. The Glow low-level graph will not replace the machine learning high-level graph, in the same way that the low-level intermediate representation in compilers does not replace the abstract syntax tree.

We aim to provide a useful compiler toolkit that will allow hardware developers to focus on implementing efficient acceleration hardware, each of which likely differ in capabilities, and use Glow for automating compilation tasks such as instruction selection, memory allocation and graph scheduling. The full compiler toolkit is open-source and publicly available¹.

¹<http://github.com/pytorch/glow>

Cadence®, Esperanto®, Habana®, Intel®, and Qualcomm Technologies Inc.®, a subsidiary of Qualcomm Incorporated®, have committed to supporting Glow in future silicon products. Each of their accelerators will likely differ in capabilities, and can use Glow for automating compilation tasks such as instruction selection, memory allocation, and graph scheduling.

2 Related Work

2.1 Relationship to Neural Network Frameworks

Frameworks such as PyTorch [4], Caffe [5], and TensorFlow [6] have found success by providing a useful way for developers to create neural network models, and executing them on specific architectures. However, supporting new architectures and operators is not scalable, because adding a new operator requires it to be implemented on each supported architecture, and adding a new architecture requires all operators be implemented for it. Glow is designed to consume a neural network compute graph, optimize it, and code generate for it for a diverse set of backends in a more scalable way. This includes target-independent optimizations and analysis prior to efficiently targeting a specific backend.

ONNX [7] is an open-source format for representing and serializing AI models. It allows for interoperability between different AI frameworks, allowing compute graphs from one framework such as PyTorch to be converted to or from another framework such as Cognitive Toolkit (CNTK) [8].

2.2 Compiler-Related Projects

Several prior systems use a compiler-oriented approach to optimizing neural networks. TensorFlow’s XLA [2] compiles neural networks for CPUs, GPUs and accelerators. It is a practical compiler that solves actual problems. XLA is used in production to drive a massive fleet of accelerators at Google. It lowers nodes into primitive linear algebra operations, and then calls into a backend-specific library for different backend (such as Eigen [9] for CPUs, or cuDNN [10] for GPUs) to perform the bulk of computation. We point out that XLA emits vectorized LLVM intermediate representation (IR) [11] for some nodes (such as dot), and relies on the LLVM vectorizer [12] for other nodes. It aims to provide a backend flexibility

```
for (...) A[i] = 3;
for (...) A[i] = 4;
return A[0];
```

Figure 1: Compilers struggle to analyze and optimize this code when the two loops come from two different nodes in the dataflow graph.

for TensorFlow, in a similar way that Glow is working toward providing for PyTorch and other neural network frameworks.

TVM/NNVM [3, 13] lowers nodes into a low-level Halide-based IR wherein loop-based optimizations can be performed. Halide [14] is then used to generate LLVM or CUDA/Metal/OpenCL source code. On the other hand, DLVM [15] lowers DLVM IR into LLVM IR, benefiting from the LLVM’s mature optimizer and code generator.

Yet another approach is taken by nGraph [16], which consumes a framework’s (such as Tensorflow) compute graph to represent internally in a single level IR, and then lowers that to different backends such as cuDNN and MKL-DNN [17].

Finally, Tensor Comprehensions [18] provides a language for neural network developers to specify their networks such that a JIT compiler can algorithmically search for the most efficient execution plan possible. This execution plan is then generated in a language suited for a specific backend, such as CUDA for a GPU, and compiled by a compiler for that language. Tensor Comprehensions is a good solution for programmers that seek to create new operators that do not exist today and execute them efficiently.

Similar to Glow, these systems include one or more levels of IR (Section 3) which represent the compute graph of some neural network model. Additionally, like Glow, many represent tensors as first-class members with a shape and an element type. Glow uses multiple levels of its own IR in the entire stack, and leaves it up to each backend to implement further lowering if desired. For example, Glow’s CPU backend executes low-level Glow instructions and calls into its own libjit standard library kernels implemented in C++ and compiled with LLVM (Section 5).

3 Intermediate Representation

3.1 Motivation

In this section we describe the motivation for having a high-level intermediate representation (IR). Neu-

ral networks are dataflow graphs where each operation represents some mathematical operation, such as element-wise add or matrix multiplication. One way to compile this graph into an executable would be to translate each mathematical operation directly into some low-level compiler IR that contains loops and other low-level instructions. However, we believe that the high-level domain specific intermediate representation is necessary for optimizing the graph.

Consider the code in Figure 1. Two for-loops write into some memory region, and later the return-statement reads from some element in the array. Neither GCC nor LLVM were able to remove the first loop, which is redundant, or replace the load operation with the constant value ‘4’. The reason is that analyzing loops and memory is difficult. The compiler needs to prove that the indices in the loops do not overflow, and that pointers in the program do not alias, and that the result of the computation is accurate and conform with the specification of the C programming language.

Thus, trusting a sufficiently good C++ compiler to optimize neural networks is not a viable strategy, because for example it is difficult to reverse engineer a sequence of 7 loops into a convolution. Instead we have implemented a high-level intermediate representation that allows a compiler to reason about and optimize high-level constructs such as tensors and operations.

Glow is a retargetable compiler that supports a number of different backends. This means that the first few phases of the compiler are target-independent, but as you get closer to instruction selection the IR becomes more target-specific. This design is not unique to Glow. Many compilers and virtual machines use similar techniques to gradually canonicalize, optimize and lower programs into instruction streams. The first two levels of IR are shared between all compilation targets. Compiler backends may implement additional levels of intermediate representations.

3.2 High-Level IR

The high-level IR is a dataflow node-based graph representation that is similar to the graph that you may find inside Caffe. When we load a neural network model we construct this graph with a direct translation of one operator to one or more nodes. The high-level IR is a simple graph that allows basic transformations, such as replacing all uses of some node with another node, or modifying the content of in-

put tensors known at compile time (e.g. pre-trained weights). The graph is strongly typed, which means that inputs and output have a known tensor type (consisting of the tensor’s shape and element type), and that the types of nodes are verified by the compiler. For example, the element-wise add instruction must operate on operands of the same type.

Some strongly-typed programming languages represent dynamic types at runtime in a safe way. Swift [19] generics are an example for such type system that allows compilation for unknown yet constrained types. We have considered the idea of developing some kind of parametric tensor types to support features such as varying batch sizes. However, we have decided to implement a simple strict type system instead and let the high-level machine learning framework specialize the computation before constructing the Glow graph. We evaluated the mechanisms that the modern programming languages use to implement generics and concluded that most hardware accelerators do not support some of these mechanisms. Production systems that use Glow may generate multiple Glow graphs for different batch sizes, or recompute the graph just-in-time.

The Glow graph is structured as a module that contains multiple functions that contain multiple nodes. Storage nodes, which are similar to global variables in C programs, are owned by the module and accessible to all functions of the same module. All other nodes are owned by functions and represent the different operations of a neural network. For example, Convolution, MaxPool, MatrixMultiply, and so on are represented as nodes. These nodes are able to reference and access storage nodes that are owned by their containing module.

Storage nodes are the base class for, and are implemented as, Constant and Placeholder nodes. Constant nodes are backed by a concrete, known tensor at compilation time. Thus, the optimizer can inspect and optimize them as it sees fit. For example, the optimizer is able to delete unused Constant nodes, transpose them, quantize them (Section 4), perform constant propagation, etc. An example of a Constant node is the pre-trained weight input to a Convolution node during inference.

Placeholder nodes are symbolic nodes that may have backing tensors assigned or changed after compilation. This means that unlike Constant nodes, the optimizer cannot inspect or optimize the contents of Placeholder nodes. If the same function is compiled using different backing tensors bound to the

Placeholder nodes, the semantics of the program are unchanged. Inputs and outputs of Glow programs should be modeled using Placeholder nodes. An example of a Placeholder node is the input image data tensor for image classification in a convolutional neural network; this input image data tensor can be changed without recompiling the function.

As an end-to-end example, one module could contain both an inference function and the gradient of that inference function. The weights (to be trained by the gradient function) would be created as Placeholder nodes, as they are not constant during training. Executing the gradient function would update the tensors backing the weight Placeholder nodes. The nodes in the inference function could then reference and access the weight Placeholder nodes (backed by the now-trained weights tensors). Thus, the Placeholder nodes can be converted to Constant nodes, and the inference function can be better optimized during compilation, knowing these nodes are Constant.

The compiler has a debug method for dumping textual and graphical representations of the graph. Figure 2 depicts the compute graph of a regression of A automatically differentiated by Glow, with the value of Placeholder node A updated with the gradient of the expression. Glow lowers the nodes that compute the gradient of the expression and the stochastic gradient descent (SGD) node into a sequence of low-level operators (Sub, Mul, Add, and Save). The different compiler backends do not need to implement support for the DivGrad or SGD nodes.

By contrast, classic machine learning frameworks that are not able to automatically generate fused kernels (Section 5.2) need to implement hundreds of CUDA and CPU compute kernels that represent the un-lowered operators. This limits their ability to support new kinds of hardware and ties them to one or two major hardware vendors.

3.3 Predication

Predication is a well-known technique to control the execution of some node or instruction by means of a boolean flag. If the value of the flag at runtime is set to ‘false’ then the predicated node or instructions may return any value. A correct program should know to ignore the output of the predicated instruction because it could be zeros or uninitialized memory. The type of the flag must be a boolean value or a vector of booleans that matches the batch size. Predicates could accelerate the performance of some networks by

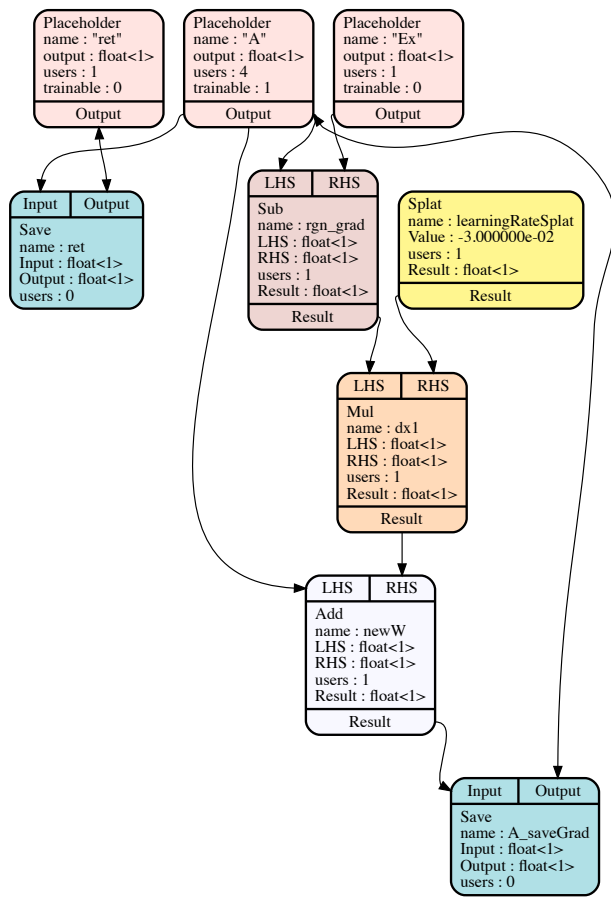


Figure 2: A lowered compute graph in Glow’s high-level IR, representing a regression of A , automatically differentiated by Glow.

avoiding some computation. It can particularly be useful when applied to Recurrent Neural Networks [20], because different elements of the batch may have different lengths and do not need to perform the same amount of computation.

3.4 Node Lowering

The Glow compilation pipeline solves the problem of targeting a large number of opcodes to many different targets. Modern machine learning frameworks support hundreds of operators on many different hardware backends. The approach that is taken by classic machine learning frameworks is to implement each opcode for each hardware target. In such frameworks, ReLU would be implemented once for the GPU, once for the CPU, once for some mobile DSP accelerator, and so on. This approach does not scale as the number of opcodes and the number of hardware targets

increase.

Instead, Glow takes a different approach. Instead of compiling the high-level operators directly, Glow performs "node lowering". In this phase, the compiler breaks the high-level operator nodes into low-level linear algebra operator nodes. For example, the Fully-Connected layer is represented as a matrix multiplication followed by broadcasted add. Different compiler backends do not have to implement the FullyConnected layer and a dozen other high-level opcodes, just the low-level matrix multiplication.

This lowering phase drives many of the design decisions of the compiler. In Glow, lowering is performed as part of the high-level graph as described above, prior to moving to low-level IR (Section 3.5). This is due to a number of reasons. First, the new lowered graph may allow for additional graph-level optimizations. Second, the new graph structure may affect the decisions of the instruction scheduler. And third, after lowering we allow the backends to perform additional target-specific optimizations on the lowered graph.

The lowering phase comes after the graph is differentiated. Because the lowering transformation does not preserve the semantics of the graph, it is not possible to differentiate the graph for certain operators. For example, the Regression node (which produces gradient when optimizing total squared error) becomes a no-op for the inference case, but is translated into an element-wise subtract for the training case. Performing the lowering before differentiation would prevent us from performing the correct lowering of the Regression node.

3.5 Low-Level IR

After optimizing the graph with target-independent optimizations, and lowering from high-level operator nodes to linear algebra operator nodes, the code is further lowered into the low-level IR in a phase that is called "IRGen" (which stands for IR generation)². This is a one-to-many translation where each high-level node is translated into one or more instructions.

The low-level IR enables a different kind of target independent optimizations that are not possible with the high-level graph format. This is an instruction-based representation that operates on tensors that are referenced by address. This gives the compiler the ability to perform low-level memory optimizations

²IRGen is optional; backends can do it if they have their own software stack that prefers to consume the Node representation of the program.

that are not possible at the high-level, because memory is not represented directly. An example of such a transformation is the optimization that allows certain operations to transform some buffers in-place, such as element-wise arithmetic.

In the context of hardware acceleration, the low-level instruction-based representation allows the compiler to represent device-specific operations such as asynchronous DMA operations. Hiding the latency of memory operations is important for utilizing the execution units of the hardware effectively, and the instruction-based representation allows the compiler to create a schedule that hides the latency of the memory operations.

The IR is strongly typed and each instruction operand kind has known parameter types. It is designed to be used as an in-memory form, though can be dumped to human readable assembly-like format.

A function in IR form contains two sections: 'declare' and 'program'. In the first section of the IR we declare a number of memory regions that live throughout the lifetime of the program. This is similar to global variables in C. The second part of the IR is a list of instructions. Each variable is annotated with the kind of initialization that the program should do.

There are two kinds of memory regions which correspond to these two sections: global memory regions (found in 'declare') and locally allocated regions (found in 'program'). The locally allocated memory regions are similar to 'alloca' in LLVM IR³. Memory regions are strongly typed, which means that the kind of type of tensor that the region represents is known.

Instructions operate on either these global memory regions or locally allocated regions. Each operand is annotated with one of the qualifiers '@in'/'@out'/'@inout'. '@in' means that the buffer is read from. '@out' means that the buffer is written into. And '@inout' means that the instruction may read and write into the buffer. These operand qualifiers help the optimizer decide when it is legal to perform certain optimizations, such as copy elimination or buffer sharing. Instructions may have other attributes that specify the legality of some optimizations. For example, some instructions require that the data from the forward pass would be kept around for the backward pass, so if the program is not optimized for inference-only mode then certain memory optimizations cannot happen.

³<http://llvm.org/docs/LangRef.html#alloca-instruction>

```

declare {
  %input = weight float<8 x 28 x 28 x 1>,
    broadcast, 0.0
  %filter = weight float<16 x 5 x 5 x 1>,
    xavier, 25.0
  %filter0 = weight float<16>, broadcast,
    0.100
  %weights = weight float<10 x 144>, xavier,
    144.0
  %bias = weight float<10>, broadcast, 0.100
  %selected = weight index<8 x 1>
  ...
  %result = weight float<8 x 10>
}

program {
  %allo = alloc float<8 x 28 x 28 x 16>
  %conv = convolution [5 1 2 16] @out %allo,
    @in %input, @in %filter3, @in %bias0
  %allo0 = alloc float<8 x 28 x 28 x 16>
  %relu = max0 @out %allo0, @in %allo
  %allo1 = alloc index<8 x 9 x 9 x 16 x 2>
  %allo2 = alloc float<8 x 9 x 9 x 16>
  %pool = pool max [3 3 0] @out %allo2, @in
    %allo0, @inout %allo1
  ...
  %deal6 = dealloc @out %allo6
  %deal7 = dealloc @out %allo7
  %deal8 = dealloc @out %allo8
  %deal9 = dealloc @out %allo9
}

```

Figure 3: Unoptimized low-level Glow IR.

Figure 3 shows an example of unoptimized Glow IR. Note that the ‘alloc’ instruction does not allocate memory; it just marks the lifetime of the activation. The low-level memory allocator is responsible for allocating all of the buffers into a single coalesced region.

3.6 Summary: The Lifetime of a Glow Instruction

This section summarizes how instructions travel from the beginning of the compilation pipeline, and through the different levels of IR and to the backends. This is a high-level overview of the compilation process:

1. The graph is either loaded via the graph loader (from ONNX or Caffe2 format), or constructed via the C++ interface.
2. The graph is differentiated if needed.
3. The graph is optimized.
4. Linear algebra node lowering takes place.
5. Additional rounds of optimizations occur, both target independent and target specific.

```

BB.newInstr("AvgPool")
.addOperand("Dest", OperandKind::Out)
.addOperand("Src", OperandKind::In)
.addMember(MemberType::VectorUnsigned, "Kernels")
.addMember(MemberType::VectorUnsigned, "Strides")
.addMember(MemberType::VectorUnsigned, "Pads")
.autoIRGen()
.autoVerify(VerifyKind::SameElementType,
  {"Dest", "Src"})
.addGradientInstr({"Dest"}, {"Dest", "Src"});

```

Figure 4: Example class-gen for the Average Pool instruction.

6. The graph is scheduled into a linear sequence of nodes that minimizes memory usage.
7. IRGen converts the low-level graph into instructions.
8. Low-level IR optimizations are performed.
9. Backend-specific optimizations and code generation are performed.

3.7 ClassGen

Glow uses automatic code generation techniques (class-gen) for defining instructions and nodes. The purpose of the automatic code generation tools in Glow is similar to the motivation behind LLVM’s TableGen, which is to help a human develop and maintain records of domain-specific information.

The current system is capable of generating two kinds of classes: Nodes for the high-level IR and Instructions for the low-level IR. Figure 4 shows an example of the code for generating the AvgPool instruction. ClassGen generates most of the methods that instructions need to have, such as instruction equality and hashing, cloning, printing, verification, etc.

4 Quantization

In the context of machine learning, quantization is the process of converting a neural network from floating-point to integer arithmetic. Arithmetic using small integers is more efficient than the computation of full-width floating-point numbers, and additionally decreases memory usage.

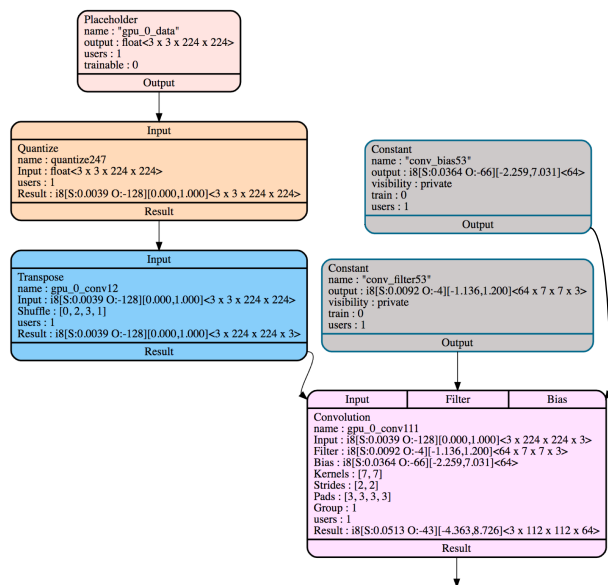
Glow is able to convert floating-point-based networks into signed 8-bit integer networks. The canonical quantization representation is using signed integers, though it is possible to support other quantization formats. Glow uses profile-guided quantization,

4.1 Tensor Representation

```
value = (input - offset) * scale
```

4.2 Profile-Guided Quantization

We use profile-guided information to estimate the possible numeric range for each stage of the neural network. Our quantization conversion works using a two-phase process. First, we statically instrument the network with special profiling nodes that record the ranges of activations that flow in the network, optimize the network including these profiling nodes, and then run inference. Then, we recompile the network using this profile information to convert the network into a quantized form, allowing for static optimization of the quantized graph. We convert portions of the network into islands of integer computation and aim to generate outputs in the range that the original floating-point network produces. Figure 5 shows a quantized subgraph from Resnet50.



4.3 Compiler Optimizations For Quantization

First, we attempt to minimize the number of conversions between floating-point tensors and integer tensors, in both directions. Some operations, such as ‘transpose’ and ‘concat’ operate on both types, and changing the representation can minimize conversions.

Second, the neural network contains ‘rescale’ nodes that change the range of the integers. These nodes are required to convert between numeric ranges that mimic the original floating-point network. However, in many cases, it is possible to fold the rescale operations into numeric-producing operations, and eliminate them.

Third, it’s possible to rescale the values in the network in order to allow fast hardware implementations of the quantized operations. For example, consider the ‘max’ operations. By converting both sides of the ‘max’ into the same scale we allow the hardware to perform a simple comparison. By normalizing both sides of the ‘max’ operation to the same scale we enable this efficient optimization.

5 CPU Backend

This section describes the implementation of the CPU backend. The Glow CPU backend compiles the low-level intermediate representation into an optimized stream of instructions. It uses LLVM to optimize and emit machine code and was tested on x86 and ARM64. The backend can emit a stand-alone object file to disk or execute code in just-in-time mode. The backend emits debug information, which makes it possible to debug Glow in a debugger and place a breakpoint in specific operator, or understand the performance of networks using a profiler.

5.1 Standard Library

One interesting aspect of the Glow CPU backend is the use of a small target independent standard library. The CPU backend needs to generate code for machine learning operators such as Convolution and SoftMax. One possibility is to call into some external library such as Eigen. This is easy to do, and many machine learning frameworks use this technique. The disadvantage of this technique is that the external binary library has no information about the specific operation that is being compiled. Some of the parameters that an optimized implementation may care about are the specific tensor sizes, the exact addresses of buffers in memory, and whether some pointer aliases another pointer.

Glow compiles a small standard library that ships with the compiler into LLVM bitcode. During compilation, Glow loads the bitcode from disk and specializes the operator implementations for the specific context. Glow replaces function arguments that represent the dimensions of some tensor or buffer addresses with constants that LLVM can optimize to generate efficient code. The compiler can decide on the kind and level of operator specialization to perform, trading compile time and binary size for performance.

Most operators are very simple and the LLVM vectorizer [12] is able to generate very efficient code. Notice that by providing the exact tensor dimensions and loop trip count the vectorizer is able to generate efficient code that does not contain pre-header legality check and scalar loop to handle the remainder odd iterations. The convolution and matrix multiplication operations are hand-optimized in C++ using the clang extended OpenCL vector syntax, and LLVM does a good job allocating registers and encoding the instructions, removing the need to use inline assembly.

```
Filter layout before transformation:
[depth, filter_x, filter_y, channel]
Filter layout after transformation:
[depth/N, filter_x, filter_y, channel, N]
```

Figure 6: Transformation of a convolution’s filter’s memory layout to optimize for SIMD memory accesses. Depth refers to the output depth of the filter, and channel refers to the input channel.

5.2 Operator Stacking

One important optimization that the CPU backend implements is stacking of data-parallel operators. Consider a sequence of operators that operate one element at a time, for example a ReLU, Add, Sub. Iterating over a large buffer multiple times is inefficient because it requires the CPU to load the memory multiple times, each time invalidating the whole cache. Instead, Glow stacks operators and performs a few data-parallel operators one after the other on the same memory location. Notice that as described above, this is not an optimization that LLVM can perform by itself and it requires a special high-level data structure.

Operator stacking is similar to operator fusion. However, when fusing multiple operators (e.g. Conv and ReLU fused together), all backends that want to support this fused operator must implement a specific kernel for each permutation of operators. In contrast, Glow’s stacking automatically creates such kernels; all of the possible permutations of data-parallel nodes are automatically fused into a fast kernel.

The approach of stacking multiple operations has many advantages. First, there is an immediate performance gain for places in the graph where data-parallel operators are placed one on top of the other. Second, backends do not need to implement kernels for all possible permutations of consecutive data-parallel nodes. And lastly, it allows Glow to lower high-level operators knowing that the backend can fuse them and recover the performance.

For example, Glow lowers the SGD (stochastic gradient descent) operator into a sequence of low-level primitives that include addition, subtraction, and multiplication. Lowering the SGD node into low-level primitives simplifies the design of the compiler by reducing the operator-space that the backend needs to handle. Operator stacking can also accelerate computation on GPUs by reducing the kernel launch overhead.

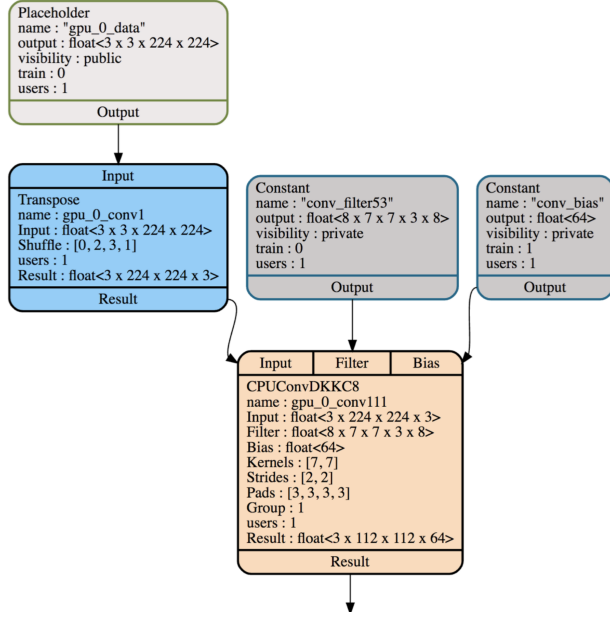


Figure 7: A subgraph from Resnet50 optimized for the CPU backend. The CPUConvDKKC8 node requires weights with a modified memory layout for efficient SIMD access (Figure 6).

5.3 Use Case: Optimizing Resnet50 for the CPU

Here we describe the way that Glow optimizes Resnet50 to generate an efficient stream of x86 instructions. Resnet50 is a residual convolutional neural network containing 54 convolutions as well as other operators such as element-wise addition, ReLU, batch normalization, max and average pooling, fully-connected, and softmax. Glow optimizes Resnet50 by performing high-level and low-level optimizations.

First, high-level transformations eliminate redundant transpose operations and merge the batch normalization operation with a convolution node. Next, the CPU backend transforms the graph into a target-specific graph that allows device-specific optimization. The CPU backend identifies three kinds of convolutions: convolutions with a small number of channels, convolutions where the size of the input activation buffer is large, and convolutions where the filter weight buffer is large. Each one of these convolutions requires a different compilation strategy. Next, the target-specific optimizer mutates the graph and generates code that matches the selected convolution. Each convolution kind uses a different filter memory layout and tile size. Figure 6 depicts the transformed filter memory layout.

```
LBB14_1:
vmovaps 3211264(%rcx,%rax,4), %ymm1
vmovaps 3211296(%rcx,%rax,4), %ymm2
vmovaps 3211328(%rcx,%rax,4), %ymm3
vaddps 6422528(%rcx,%rax,4), %ymm1, %ymm1
vaddps 6422560(%rcx,%rax,4), %ymm2, %ymm2
vmovaps 3211360(%rcx,%rax,4), %ymm4
vaddps 6422592(%rcx,%rax,4), %ymm3, %ymm3
vaddps 6422624(%rcx,%rax,4), %ymm4, %ymm4
vmaxps %ymm0, %ymm1, %ymm1
vmaxps %ymm0, %ymm2, %ymm2
vmaxps %ymm0, %ymm3, %ymm3
vmovaps %ymm1, 6422528(%rcx,%rax,4)
vmovaps %ymm2, 6422560(%rcx,%rax,4)
vmaxps %ymm0, %ymm4, %ymm1
vmovaps %ymm3, 6422592(%rcx,%rax,4)
vmovaps %ymm1, 6422624(%rcx,%rax,4)
addq $32, %rax
```

Figure 8: A loop in x86 assembly as generated by the Glow CPU Backend, with a fused element-wise addition and ReLU (max) operation.

This 5-dimensional tensor layout allows for consecutive SIMD memory access. The N parameter is selected based on the iteration order and the blocking strategy for the convolution. The CPU backend traverses the graph and replaces any convolutions it would like to optimize in this way with this specialized convolution. This can be seen in Figure 7.

The second parameter that the compiler controls is the size of the convolution tile. Glow selects a processing tile that depends on the size of the first level cache of the processor.

Next, the low-level optimizer optimizes the instruction stream by shrinking the lifetime of memory allocations for the activations, and then performs static memory allocation for the whole network into a single buffer. This reduces the mutable memory footprint of the network. From this point in the compilation pipeline the compiled code can refer to pointers in memory.

Finally, the compiler performs efficient code generation for the non-convolution parts of the network. For example, Figure 8 depicts the generated assembly for some part of the network. The compiler fused two unrelated element-wise operations into a single loop. The Add and Max operations are performed on the same memory buffer without reading the memory twice.

6 Glow Runtime

After compiling a model, Glow provides a runtime that is capable of partitioning models, queueing re-

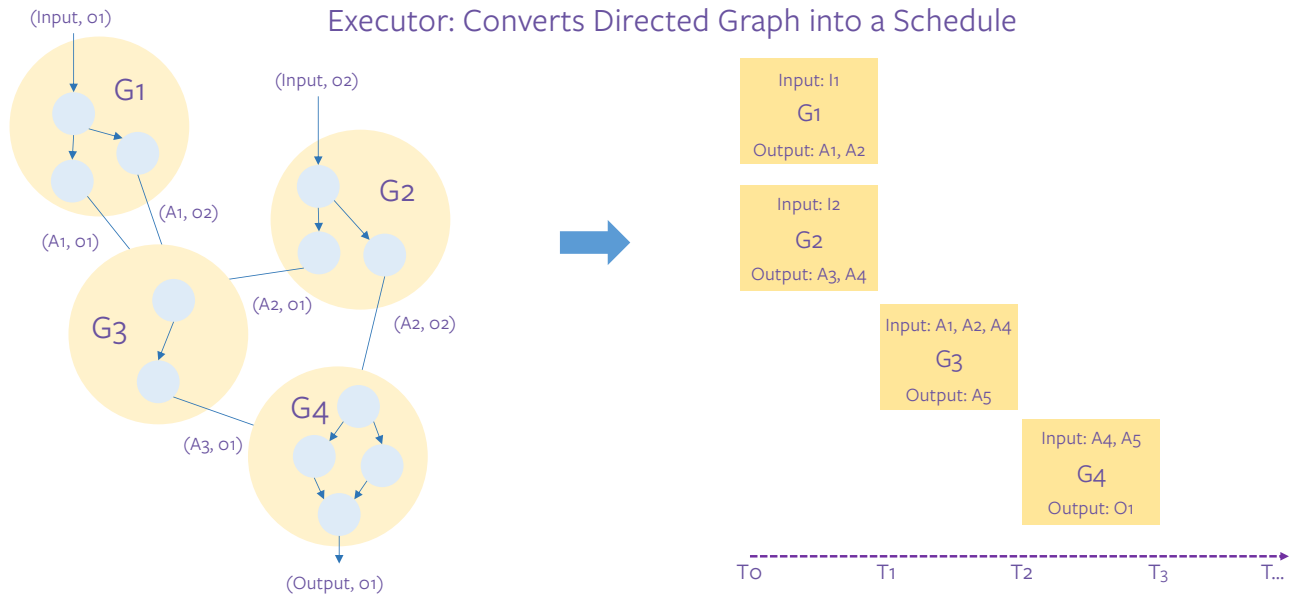


Figure 9: A simple example showing a graph partitioned into multiple sub-graphs, themselves making up a directed graph, and then converted into a schedule by the executor.




Figure 10: Four networks (A, B, C, D) have been assigned to five accelerators. A and B have been partitioned, and B has been duplicated on multiple accelerators.

quests, and executing models across multiple devices. It provides a host level abstraction for compiling and loading models and handling concurrent inference requests on all those models. The runtime allows users of Glow to target one interface without needing to worry about the underlying hardware changing, and since it manages the host’s accelerator cards, it can take advantage of its intimate knowledge of the hardware to do the best partitioning job possible.

6.1 Glow Runtime Components

The **Partitioner** splits a network into sub-networks that can be run on multiple devices. Depending on each accelerator’s available memory and the size of the weights of a model, we may want or need to partition an input network into sub-graphs across multiple accelerators in order to saturate each accelerator. A network is divided into sub-networks based on different criteria: memory constraints, estimated time cost, and communication cost between devices.

The **Provisioner** assigns partitioned sub-graphs to specific devices and calls into the backend and Device Manager to compile and load each sub-graph onto a device. An example can be seen in Figure 10.

The **Device Manager** serves as an abstraction for the physical device. It handles network loading, memory transfers, execution on the device, and tracks hardware state. Just like we have a different backend per device type, there is a DeviceManager class per device type and an instance of DeviceManager per physical accelerator on the host.

The **Executor** handles the execution of a network. It tracks each sub-network’s execution state and propagates sub-network inputs and outputs. The Executor is responsible for asynchronously handling incoming inference requests for a network and returning the collated results. Figure 9 shows a simple example of a partitioned, directed graph converted into a schedule.

6.2 Execution Flow

Adding a network:

1. The Partitioner splits the network into one or more sub-networks.
2. The Provisioner compiles each sub-network and assigns them to one or more Devices.
3. One or more DeviceManagers load the sub-networks and their weights onto its associated Device.

Handling an inference request:

1. The HostManager creates a new Execution graph with intermediate storage.
2. The Executor kicks off the first sub-network execution.
3. The DeviceManager loads inputs onto the card and begins execution. When done, it reads outputs and signals completion.
4. The Executor triggers any sub-networks with satisfied dependencies.
5. When complete, the HostManager returns outputs.

7 Evaluation

We compare the performance of Glow vs. TensorFlow-1.7 and TVM. The experiments were run on a Kaby Lake Intel® Core i7-7600U (which does not support AVX-512) running on a single CPU core at 2.80 GHz. All three frameworks were compiled to support the native architecture. We used the Keras library [21] to supply and run pre-trained models for TensorFlow and TVM.

Two popular convolutional neural networks, Resnet50 [22] and VGG19 [23], are evaluated as seen in Figure 11. TensorFlow was compiled with XLA enabled. TVM was compiled with LLVM 6.0, but without auto-tuning enabled or any specialized schedules. Presented results used a batch size of 8 on all three frameworks. Each framework’s performance (in frames per second) did not vary significantly across batch sizes of 2, 4, and 8.

As seen in Figure 11, Glow is up to 2.7x faster than TensorFlow. This is due to the fact that TensorFlow calls into Eigen which implements convolution using the classic im2col followed by matrix multiplication, while Glow compiles direct convolution (Section 5.2)



Figure 11: Glow vs. TensorFlow-1.7 and TVM on an Intel® Core i7-7600U; frames per second on a single thread.

and thus avoids im2col overhead. In addition, Glow performs shape-aware code-generation.

Additionally, Glow is up to 1.3x faster than TVM. Note that we did not use autotuning and improved schedules with TVM; we expect this would improve TVM’s performance. TVM does not use im2col like TensorFlow; similar to Glow, TVM lowers nodes into a low-level IR. This IR is Halide-based, which generates LLVM IR that is compiled for the native architecture. TVM and Glow also both generate code with efficient memory access patterns such as tiling.

8 Conclusion

This paper presented the design of Glow, a machine learning compiler for heterogeneous hardware. Glow lowers the compute graph of neural networks to multi-level strongly-typed intermediate representations, enabling analyses and optimizations appropriate for each level to efficiently and scalably target many backends. We hope our efforts will enable research in the area of machine learning acceleration.

9 Acknowledgements

In addition to the core Glow team, many fellow people at Facebook have made contributions to the project, including Andrew Adams, Michel Aoun, William Arnold, Sarah Bird, Brad Cotel, Stephen Chen, Evan Cheng, Soumith Chintala, Sy Choudhury, Chris Dewan, Utku Diril, Marat Dukhan, Oniel Duncan,

Dmytro Dzhulgakov, Peter Goldsborough, Chris Gottrath, Kim Hazelwood, Yangqing Jia, Aravind Kaliah, Daya S Khudia, Changkyu Kim, Bruce Lin, Howard Mansell, Erik Meijer, Arun Moorthy, Sam Naghshineh, Maxim Naumov, Avinash Nayak, Pieter Noordhuis, Joe Pamer, Joel Pobar, Yuri Putivsky, Lin Qiao, Vijay Rao, Martin Schatz, Alexander Sidorov, Joe Spisak, Narayanan Sundaram, Andrew Tulloch, Mor Tzur, Nicolas Vasilache, Adam Weis, and Hector Yuen.

We also would like to thank Eli Bendersky, Chris Leary, Richard Wei, and Tianqi Chen for the development and release of their work to the open source.

References

- [1] John L. Hennessy and David A. Patterson. *Computer Architecture, Sixth Edition: A Quantitative Approach*, chapter 7.2. 6th edition, 2017.
- [2] XLA: Domain-specific compiler for linear algebra to optimizes tensorflow computations. <https://www.tensorflow.org/performance/xla>.
- [3] T. Chen, T. Moreau, Z. Jiang, H. Shen, E. Yan, L. Wang, Y. Hu, L. Ceze, C. Guestrin, and A. Krishnamurthy. TVM: End-to-End Optimization Stack for Deep Learning. *ArXiv e-prints*, February 2018.
- [4] Adam Paszke, Sam Gross, Soumith Chintala, Gregory Chanan, Edward Yang, Zachary DeVito, Zeming Lin, Alban Desmaison, Luca Antiga, and Adam Lerer. Automatic differentiation in PyTorch. In *NIPS-W*, 2017.
- [5] Yangqing Jia, Evan Shelhamer, Jeff Donahue, Sergey Karayev, Jonathan Long, Ross Girshick, Sergio Guadarrama, and Trevor Darrell. Caffe: Convolutional Architecture for Fast Feature Embedding. In *Proceedings of the 22nd ACM International Conference on Multimedia*, MM '14, pages 675–678, New York, NY, USA, 2014. ACM.
- [6] Martín Abadi, Paul Barham, Jianmin Chen, Zhifeng Chen, Andy Davis, Jeffrey Dean, Matthieu Devin, Sanjay Ghemawat, Geoffrey Irving, Michael Isard, Manjunath Kudlur, Josh Levenberg, Rajat Monga, Sherry Moore, Derek G. Murray, Benoit Steiner, Paul Tucker, Vijay Vasudevan, Pete Warden, Martin Wicke, Yuan Yu, and Xiaoqiang Zheng. TensorFlow: A System for Large-scale Machine Learning. In *Proceedings of the 12th USENIX Conference on Operating Systems Design and Implementation*, OSDI'16, pages 265–283, Berkeley, CA, USA, 2016. USENIX Association.
- [7] Onnx. <https://onnx.ai/>.
- [8] Frank Seide and Amit Agarwal. Cntk: Microsoft’s open-source deep-learning toolkit. In *Proceedings of the 22nd ACM SIGKDD International Conference on Knowledge Discovery and Data Mining*, KDD '16, pages 2135–2135, New York, NY, USA, 2016. ACM.
- [9] Eigen. <http://eigen.tuxfamily.org/>.
- [10] NVIDIA cuDNN. <https://developer.nvidia.com/cudnn>.
- [11] Chris Lattner and Vikram Adve. LLVM: A Compilation Framework for Lifelong Program Analysis & Transformation. In *Proceedings of the International Symposium on Code Generation and Optimization: Feedback-directed and Runtime Optimization*, CGO '04, pages 75–, Washington, DC, USA, 2004. IEEE Computer Society.
- [12] Nadav Rotem and Arnold Schwaighofer. Vectorization in LLVM. <https://llvm.org/devmtg/2013-11/#talk10>. 2013 LLVM Developers’ Meeting.
- [13] University of Washington Paul G. Allen School of Computer Science & Engineering, Amazon Web Service AI team, and DMLC open-source community. NNVM Compiler: Open Compiler for AI Frameworks. <http://tvm-lang.org/2017/10/06/nnvm-compiler-announcement.html>.
- [14] Jonathan Ragan-Kelley, Connelly Barnes, Andrew Adams, Sylvain Paris, Frédo Durand, and Saman Amarasinghe. Halide: A Language and Compiler for Optimizing Parallelism, Locality, and Recomputation in Image Processing Pipelines. In *Proceedings of the 34th ACM SIGPLAN Conference on Programming Language Design and Implementation*, PLDI '13, pages 519–530, New York, NY, USA, 2013. ACM.
- [15] R. Wei, L. Schwartz, and V. Adve. DLVM: A modern compiler infrastructure for deep learning systems. *ArXiv e-prints*, November 2017.
- [16] S. Cyphers, A. K. Bansal, A. Bhiwandiwala, J. Bobba, M. Brookhart, A. Chakraborty, W. Constable, C. Convey, L. Cook, O. Kanawi, R. Kimball, J. Knight, N. Korovai, V. Kumar, Y. Lao, C. R. Lishka, J. Menon, J. Myers, S. Aswath Narayana, A. Procter, and T. J. Webb. Intel nGraph: An Intermediate Representation, Compiler, and Executor for Deep Learning. *ArXiv e-prints*, January 2018.
- [17] Intel MKL. <https://software.intel.com/en-us/mkl>.
- [18] N. Vasilache, O. Zinenko, T. Theodoridis, P. Goyal, Z. DeVito, W. S. Moses, S. Verdoolaege, A. Adams, and A. Cohen. Tensor Comprehensions: Framework-Agnostic High-Performance Machine Learning Abstractions. *ArXiv e-prints*, February 2018.
- [19] Swift. <https://developer.apple.com/swift/>.
- [20] Ian Goodfellow, Yoshua Bengio, and Aaron Courville. *Deep Learning*. MIT Press, 2016. <http://www.deeplearningbook.org>.
- [21] Keras: The Python Deep Learning library. <https://keras.io/>.
- [22] Kaiming He, Xiangyu Zhang, Shaoqing Ren, and Jian Sun. Deep residual learning for image recognition. In *2016 IEEE Conference on Computer Vision and Pattern Recognition, CVPR 2016, Las Vegas, NV, USA, June 27-30, 2016*, pages 770–778, 2016.
- [23] K. Simonyan and A. Zisserman. Very deep convolutional networks for large-scale image recognition. *CoRR*, abs/1409.1556, 2014.